

Peter Willendrup and Erik Knudsen DTU Physics

McStas-3.0 GPU port

Agenda

- McStas on GPU via OpenACC
(a “high-level” #pragma driven access to CUDA see <https://www.openacc.org> and <https://developer.nvidia.com/hpc-sdk>)
- How well (fast) does it work?
- Simulation flow
- What did we change?
- What does not work

McXtrace

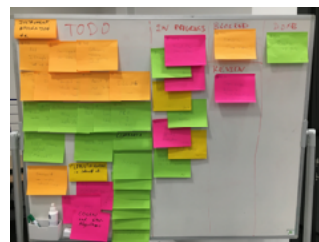
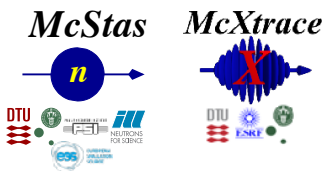


McStas



Warning:
1. Assumes previous experience with McStas
2. Does not introduce OpenACC

Main events on timeline of road toward GPU



2017: E. Farhi
initial cogen
modernisation

Fall 2018 onwards:
J. Garde further cogen
modernisation and
restructuring

October 2019 onwards:
J. Garde & P. Willendrup:
New RNG, test system, multiple
functional instruments.

January 2020:
One-week local
hackathon @ DTU
with McCode & RAMP teams

November 2020
Virtual Hackathon,
setting release scope



March 2018: Participation at
Dresden Hackathon. 1st "null"
instrument prototype runs.

October 2019:
Participation at **Espoo Hackathon**.
First meaningful data extracted.
Work on cogen and realising
we need another RNG.

November-
December 2019:
First good look at
benchmarks and
overview of what
needs doing for first
release with limited
GPU support.

2020 1st **Corona** lockdown
P. Willendrup & E. Knudsen
continue work on comp and
cogen

December 15th 2020
McStas 3.0 release!

McStas / McXtrace instrument simulation



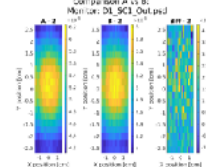
mentor: Vishal Metha

hackathon org.:
Guido Juckeland



mentor: Christian Hundt

hackathon org.:
Sebastian Von Alffthan



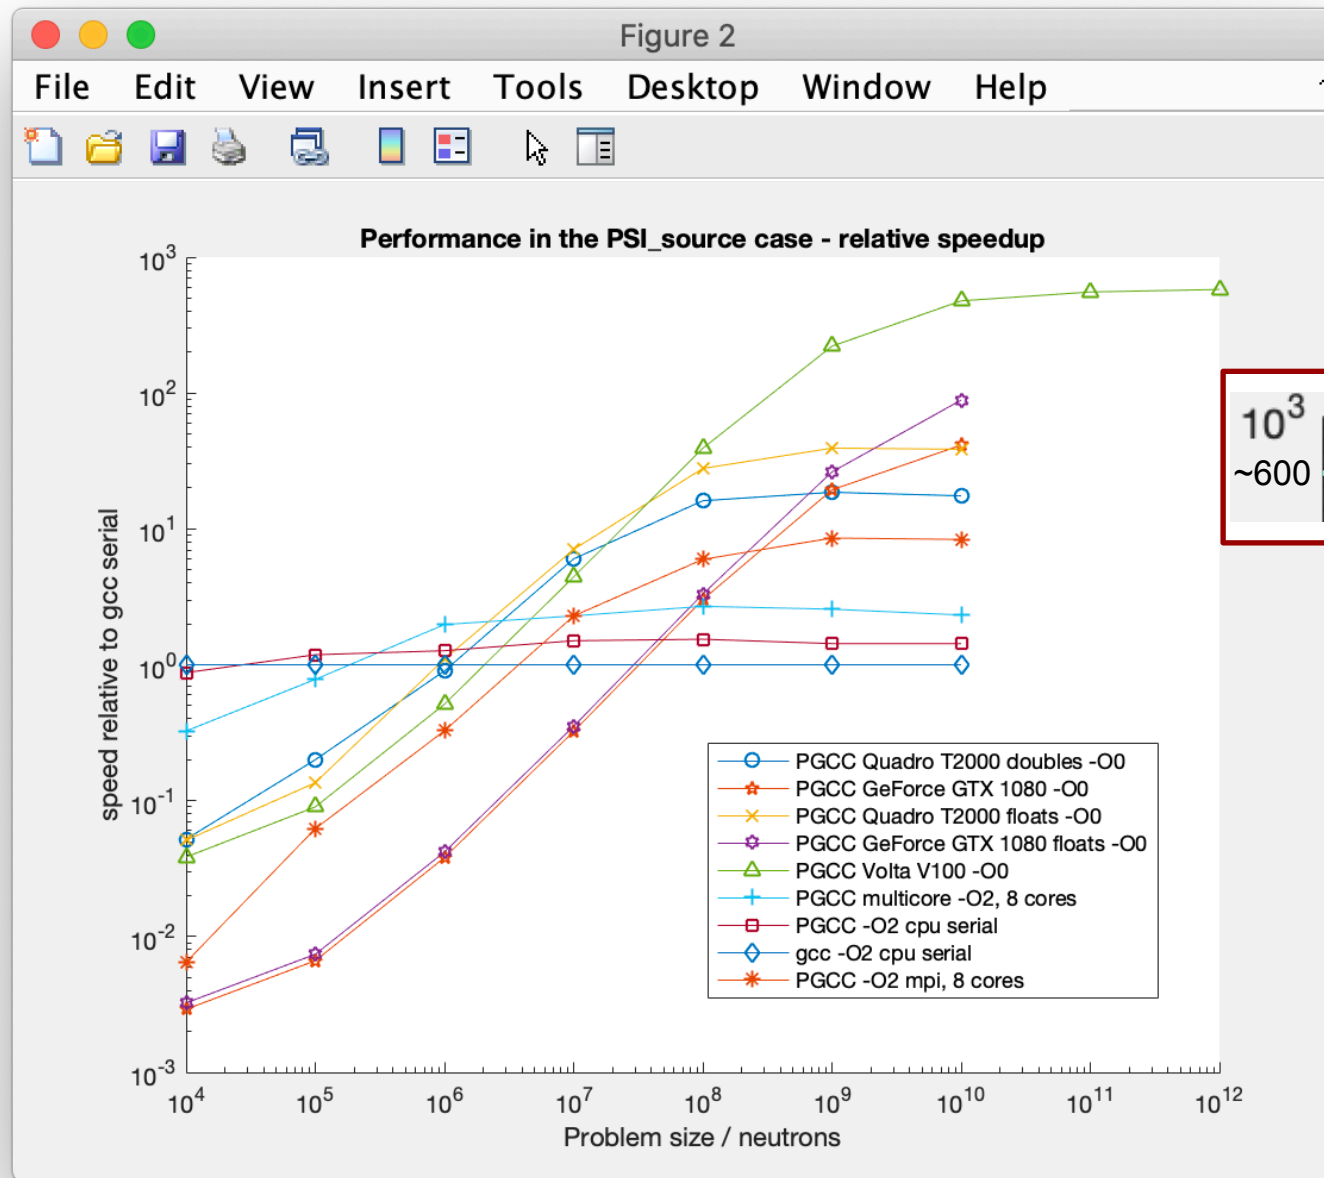
February 2020:
First release
McStas 3.0beta
with GPU
support was
released
to the public

Instrument	Configuration	Time (s)	GPU Usage (%)	Notes
BL18.1	BL18.1_1	16.12	95	Single core 56.0 s
BL18.1	BL18.1_2	5.43	5	GPU 5% usage
BL18.1	BL18.1_3

McStas heading for the GPU... November 2019 - first good look at performance.

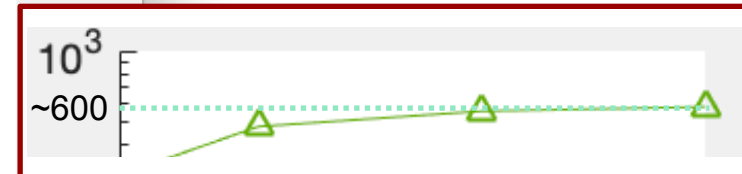
Idealised instrument
with source and monitor
only - i.e. without any
use of the ABSORB
macro.

(Likely a good indication
of maximal speedup
achievable.)



Speedup

Looks like a factor of ~600



V100 execution speedups
renormalised to wall-
clock of single-core
gcc standard simulation,

**V100 run is
600 times faster
than a single-
core CPU run**

McXtrace



McStas



McStas heading for the GPU... first benchmarking numbers from November 2019

9 instruments fully
ported, also realistic
ones like PSI_DMC

(Aug 2020: 99 instrs)

10-core MPI run,
1e9 in 200 secs



(1-core run,
1e9 would be
2000 secs)

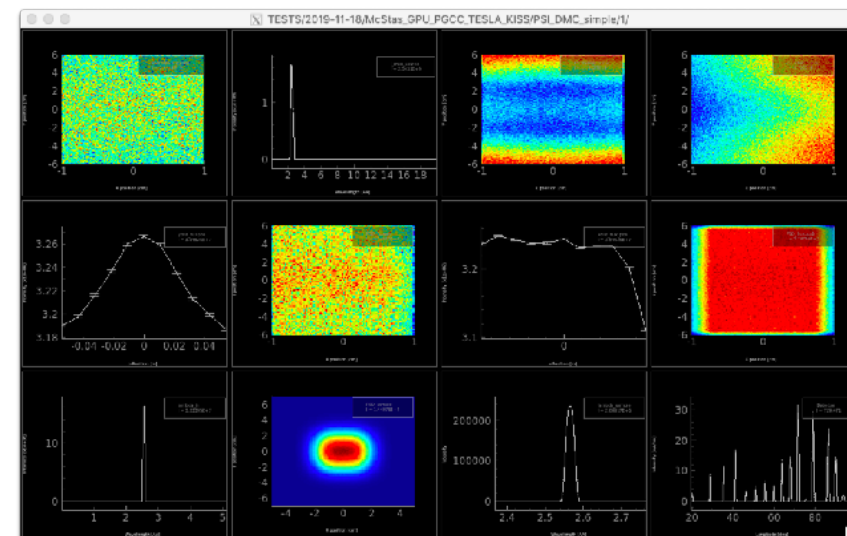
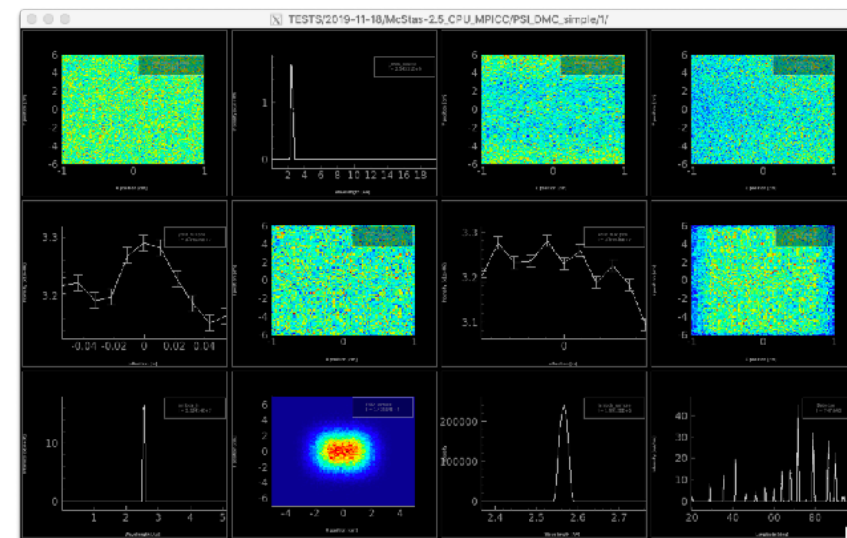
VS.

Tesla V100 run,
1e9 in 22 secs



~ i.e. 2 orders of
magnitude wrt. a single,
modern CPU core

- If problem has the right size /
complexity, GPU via OpenACC
is great!



Anatomy of a McStas GPU run (*)

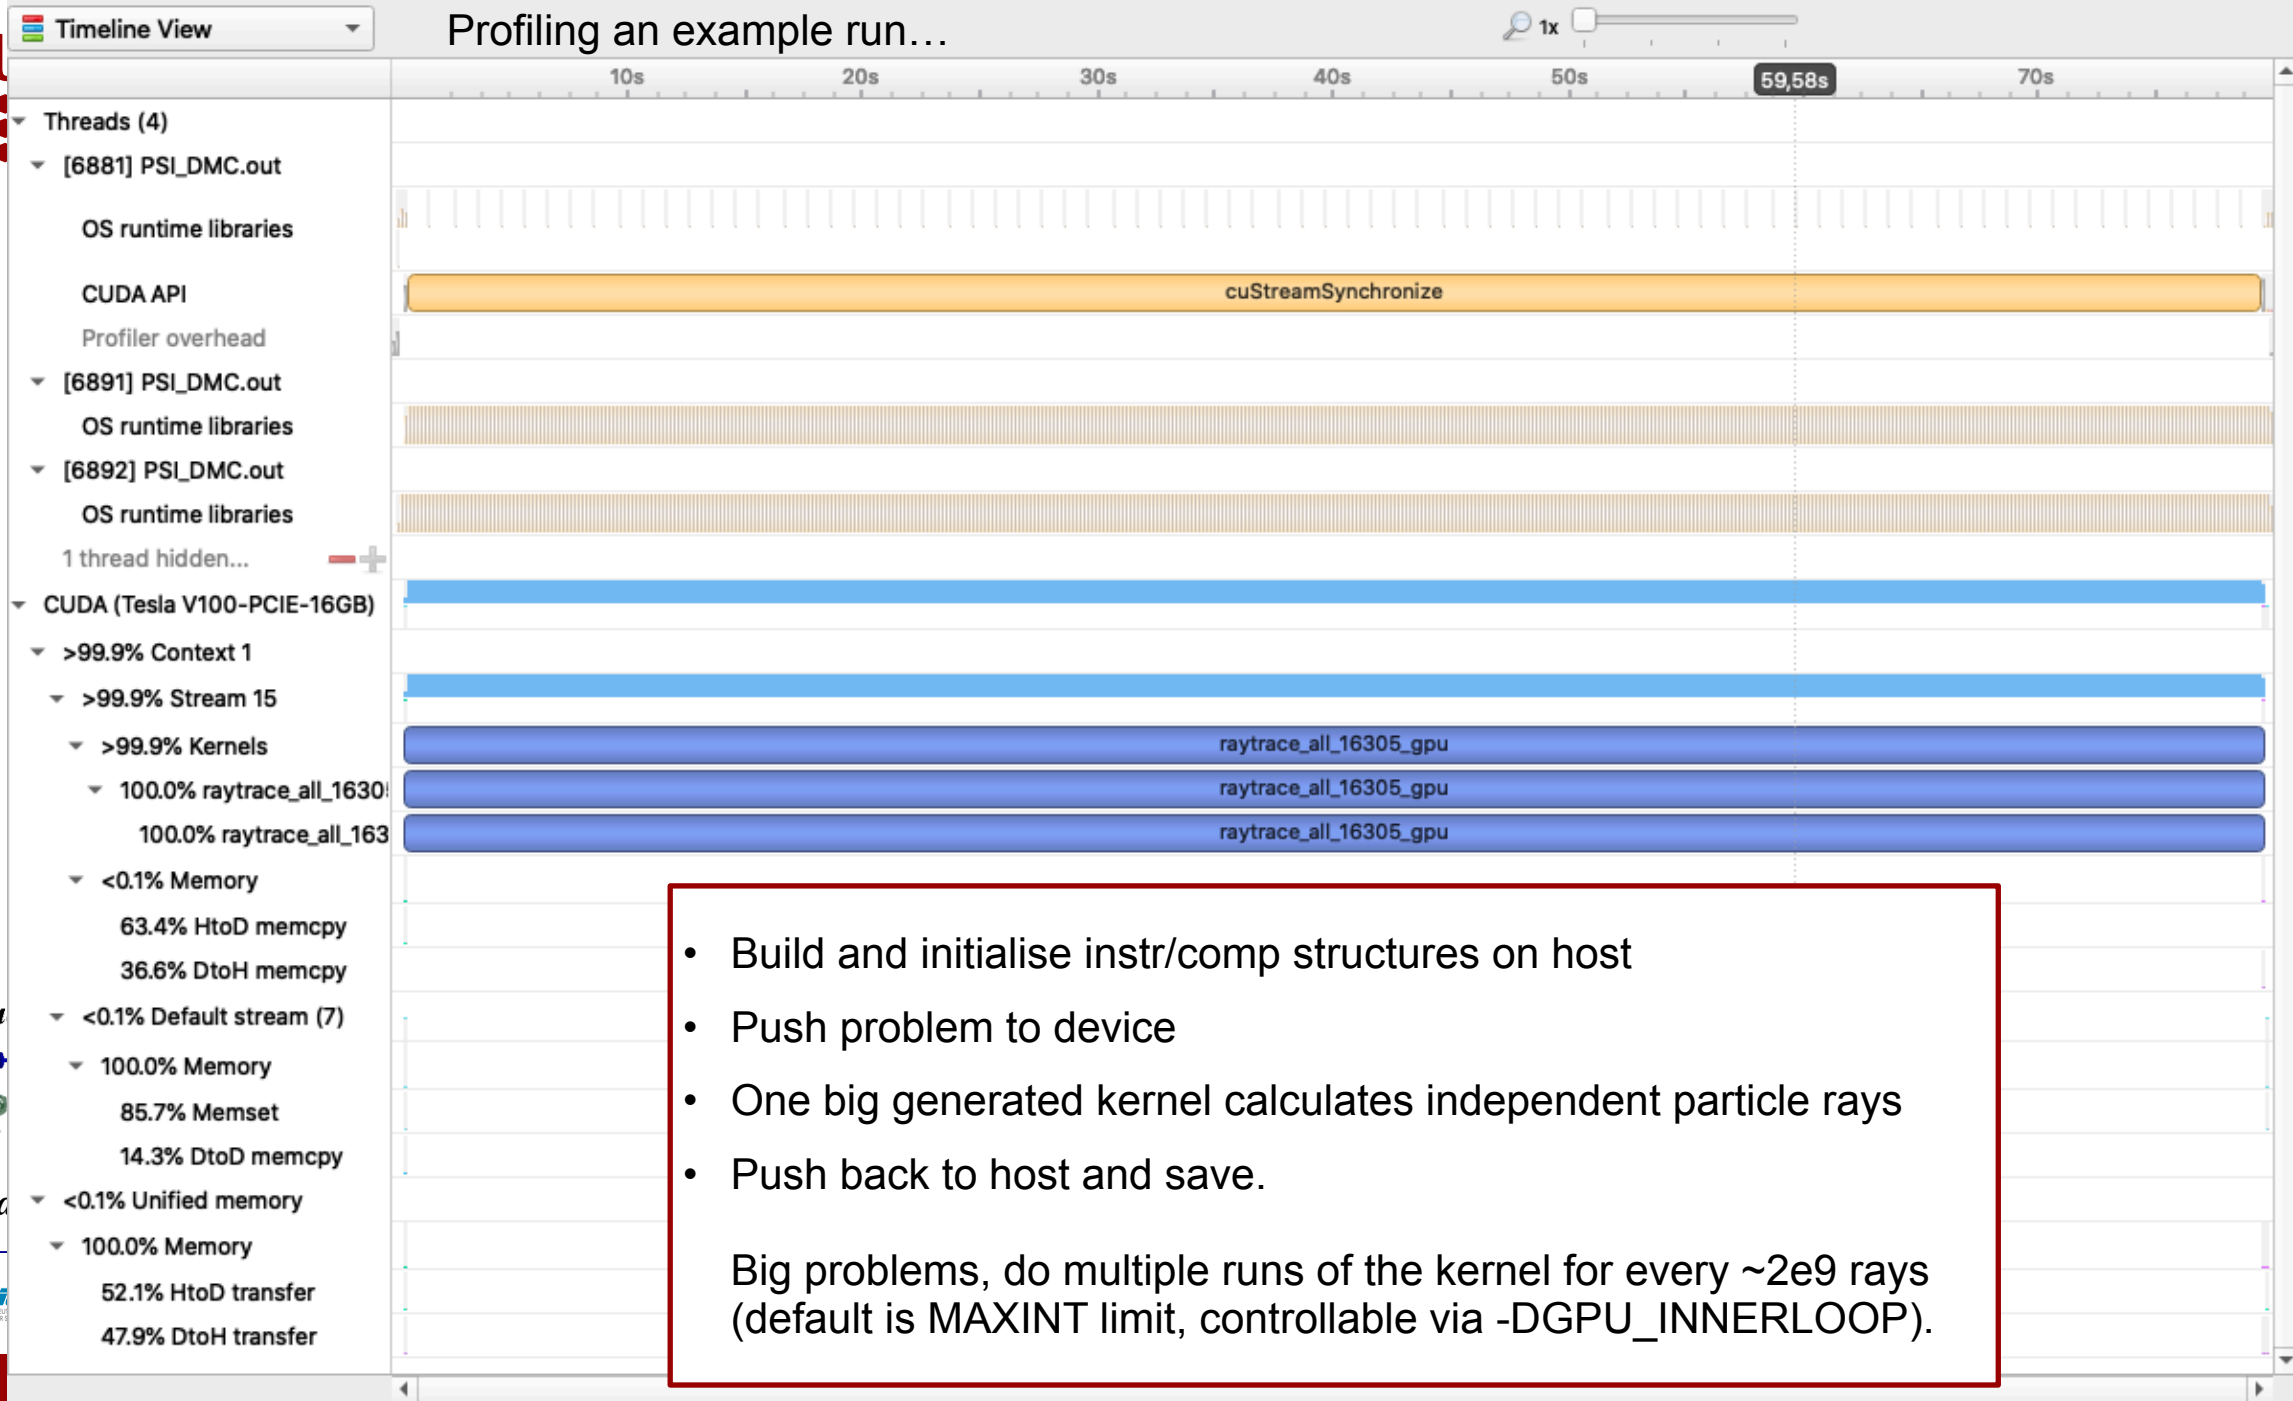
- Init, geometry, files etc. read on CPU
 - MPI if needed
- Memory-structures
 - Built on CPU
 - Marked for transfer to GPU (`#pragma acc declare create etc.`)
 - Initialised and synced across
 - Trace-loop is a `#pragma acc parallel loop`
 - Calculation performed entirely on GPU
 - Component structs (incl. e.g. monitor-arrays) synced across
- Finally and Save runs on CPU
 - MPI merge if needed



OpenACC

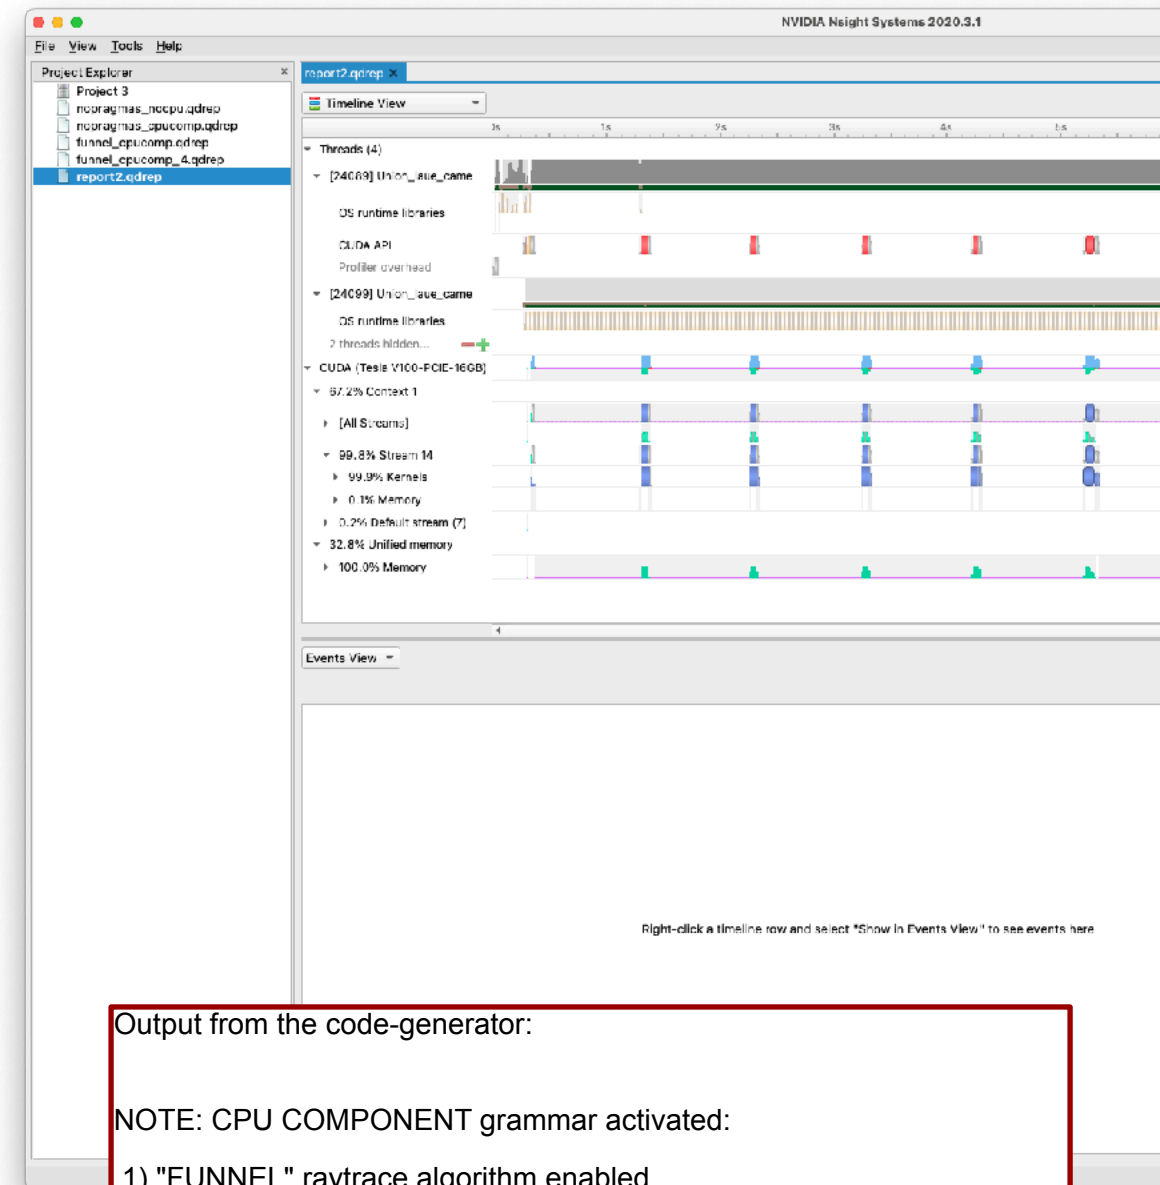
No printf's etc. available
on GPU, automatically
suppressed by `#defines`

(* Alternative layout via FUNNEL mode,
see next 2 slides)



FUNNEL mode

- Activated **explicitly** using -DFUNNEL or **implicitly** using CPUCOMPONENT in instrument or NOACC in comp header
- Has N kernels / calculation zones instead of one
 - Separation at SPLIT
 - Separation if CPUCOMPONENT in instrument file
(CPUCOMPONENT A=Comp(vars=pars...))
 - Separation if a component has NOACC in the header
(See e.g. Multilayer_sample, Union_master)
- Each of these “calculation zones” is finalised before the next one initiated.
- Example:
Union: Instrument up to Union_master can be GPU, then CPU, then GPU again
 - Can be as slow as single cpu...
 - Copying back and forth to/from GPU is costly...



Compiler settings used for GPU:

```
nvc -ta=tesla,managed -Minfo=accel -DOPENACC
```



Generate Tesla code. "compute capability" e.g. `tesla:cc70` may be specified to indicate specific card.

Use CUDA shared memory for host-device-host allocation. Needed for our 2D-arrays at present, may include penalty, we could get rid.

Give accel debug information

Main "enable GPU"/OpenACC switch

(McStas 3.0 `mcrun` is preconfigured on Linux - excluding `-Minfo=accel`, simply use `mcrun --openacc` when compiling, can also combine with e.g. `--mpi=N`)

What doesn't work



- **Function pointers are not available on GPU**
 - Solutions:
 - Code around if possible (integration routine pr. specific function to be integrated...)
 - Mark the component NOACC
- **Variadic functions are not available on GPU**
- **Anonymous structs as comp pars are not available on GPU**
 - Unfold into comp struct
- **User-defined fieldfunctions for polarisation had to be abandoned**
 - No solution yet, may become handled via grammar
- **External libs generally can not be used in TRACE** (“#pragma....” hard to add on 3rd party codes)
 - Handle in INIT / FINALLY (MCPL)
 - NOACC (GSL etc.)
- **Union master is for now NOACC**, will eventually become supported on GPU
- (Looks like we may have implemented a BUG in the NeXus/Mantid stuff...)

Conclusions

- It really does work nicely!
- **Code changes** much **less invasive** than envisioned!
- It often gives a speedup of **1-2 orders** of magnitude over 1 cpu
- **Most things work**
(we have workarounds or solutions in the pipe for the rest)
- McStas 3.0 is as of yet “ported” to GPU but **not fully “optimised” performance-wise**, we will try to go to another Hackathon
- **Union** needs a dedicated **Hackathon**



The team, Nvidia mentors and Hackathon hosts :-)



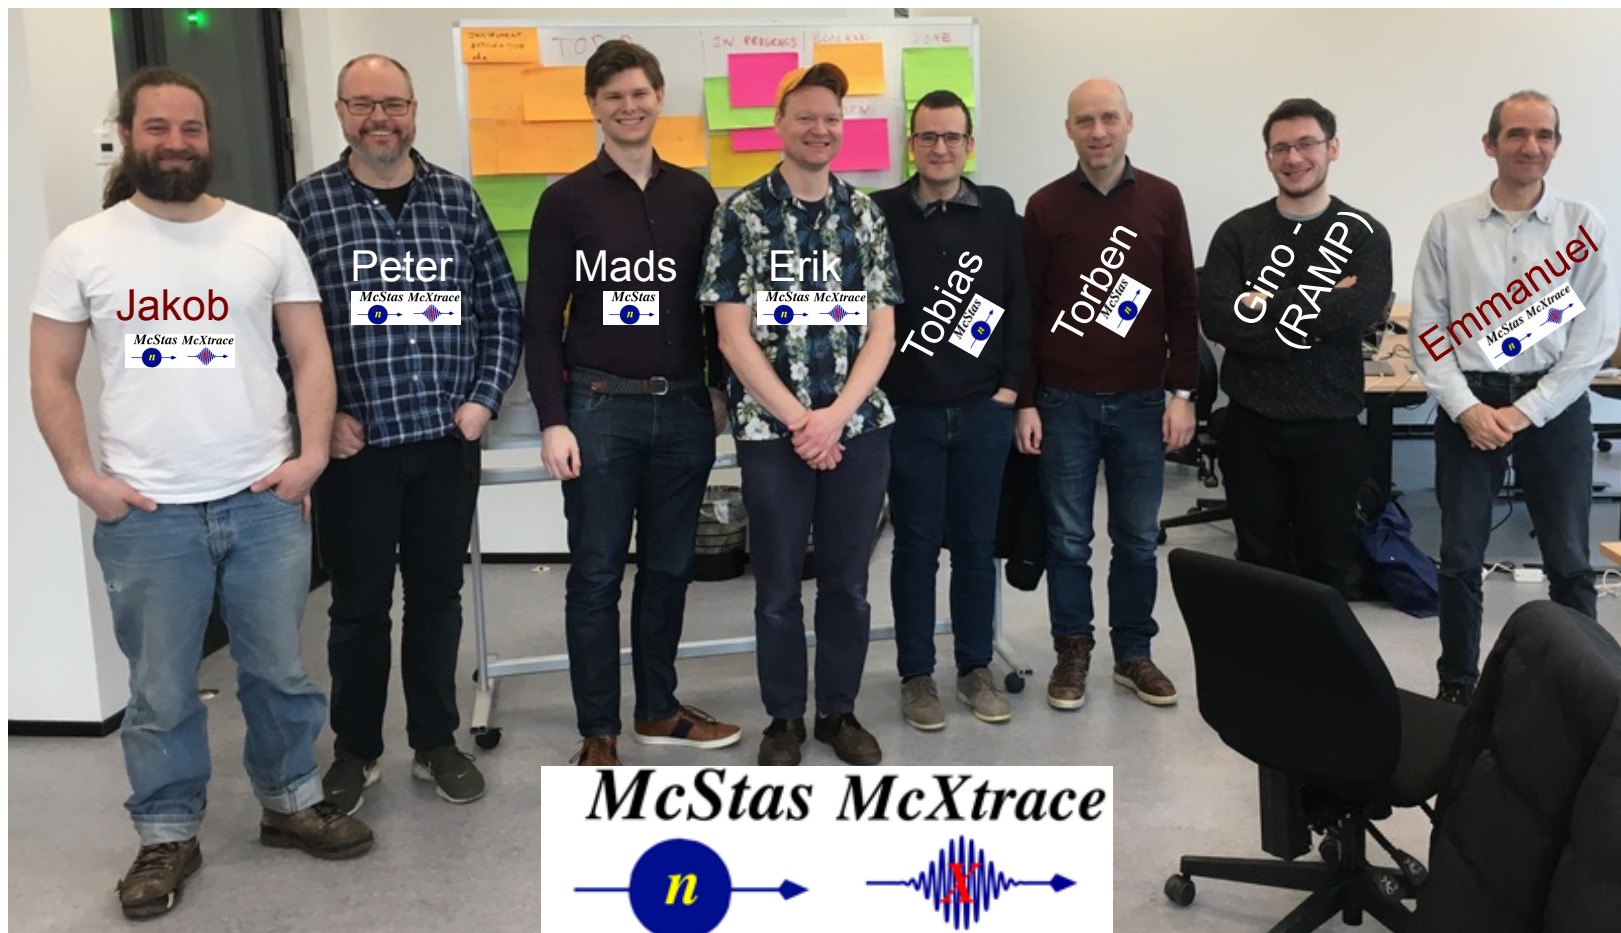
Vishal Metha



Christian Hundt



Alexey Romanenko



Guido Juckeland



Sebastian von Althaus

McXtrace



McStas

