

Peter Willendrup and Erik Knudsen DTU Physics

# **mcxtrace-1.x vs. mcxtrace-3.0, status and elements of the GPU port**

# Agenda

- McXtrace on GPU via OpenACC
  - (a “high-level” #pragma driven access to CUDA see <https://www.openacc.org> and <https://developer.nvidia.com/hpc-sdk>)
- How well (fast) does it work?
- Simulation flow
- What did we change?
- What needs doing on an instrument / component?
- What does not work

Warning:  
1. Assumes some previous experience with McXtrace  
2. Does not introduce OpenACC

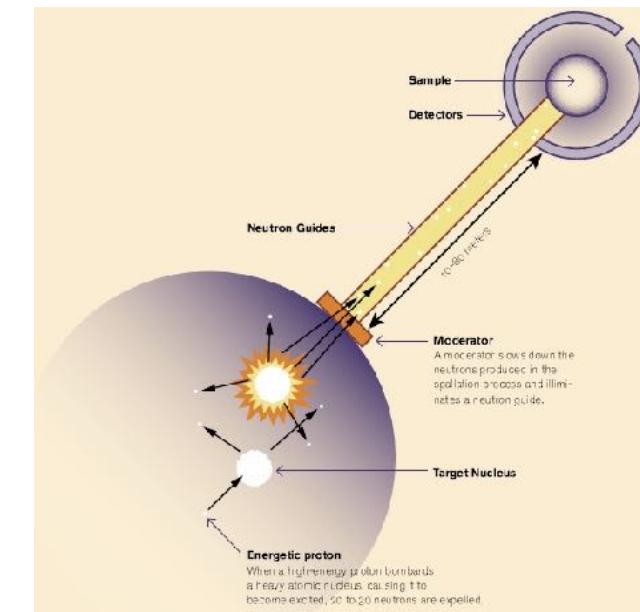
# McStas and McXtrace overview

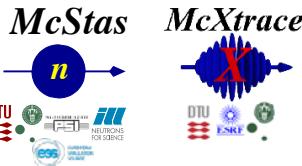
- Portable code (Unix/Linux/Mac/Windoze)
- Ran on everything from iPhone to 1000+ node cluster!



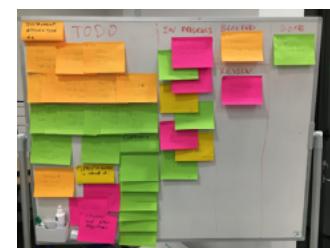
- 'Component' files (~200) inserted from library
  - Sources
  - Optics
  - Samples
  - Monitors
  - If needed, write your own comps
- DSL + ISO-C code gen.

+ NVIDIA GPU's





Fall 2018 onwards:  
J. Garde further cogen  
modernisation and  
restructuring



2017: E. Farhi  
initial cogen  
modernisation

October 2019 onwards:  
J. Garde & P. Willendrup:  
New RNG, test system, multiple  
functional instruments.

January 2020:  
One-week local  
hackathon @ **DTU**  
with McCode & RAMP teams

November 2020  
**Virtual Hackathon**,  
setting release scope

- December 15th 2020  
**McStas 3.0 release!**
- November 24th 2021  
**McStas 3.1 release!**
- February 8th 2022  
**McXtrace 3.0 release!**



March 2018: Participation at **Dresden** Hackathon. 1st "null" instrument prototype runs.

McStas / McXtrace instrument simulation



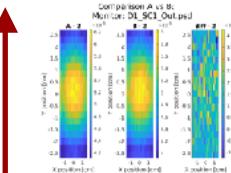
mentor: Vishal Metha

hackathon org.:  
Guido Juckeland



mentor: Christian Hundt

hackathon org.:  
Sebastian Von Alfthan



November-  
December 2019:  
First good look at  
benchmarks and  
overview of what  
needs doing for first  
release with limited  
GPU support.

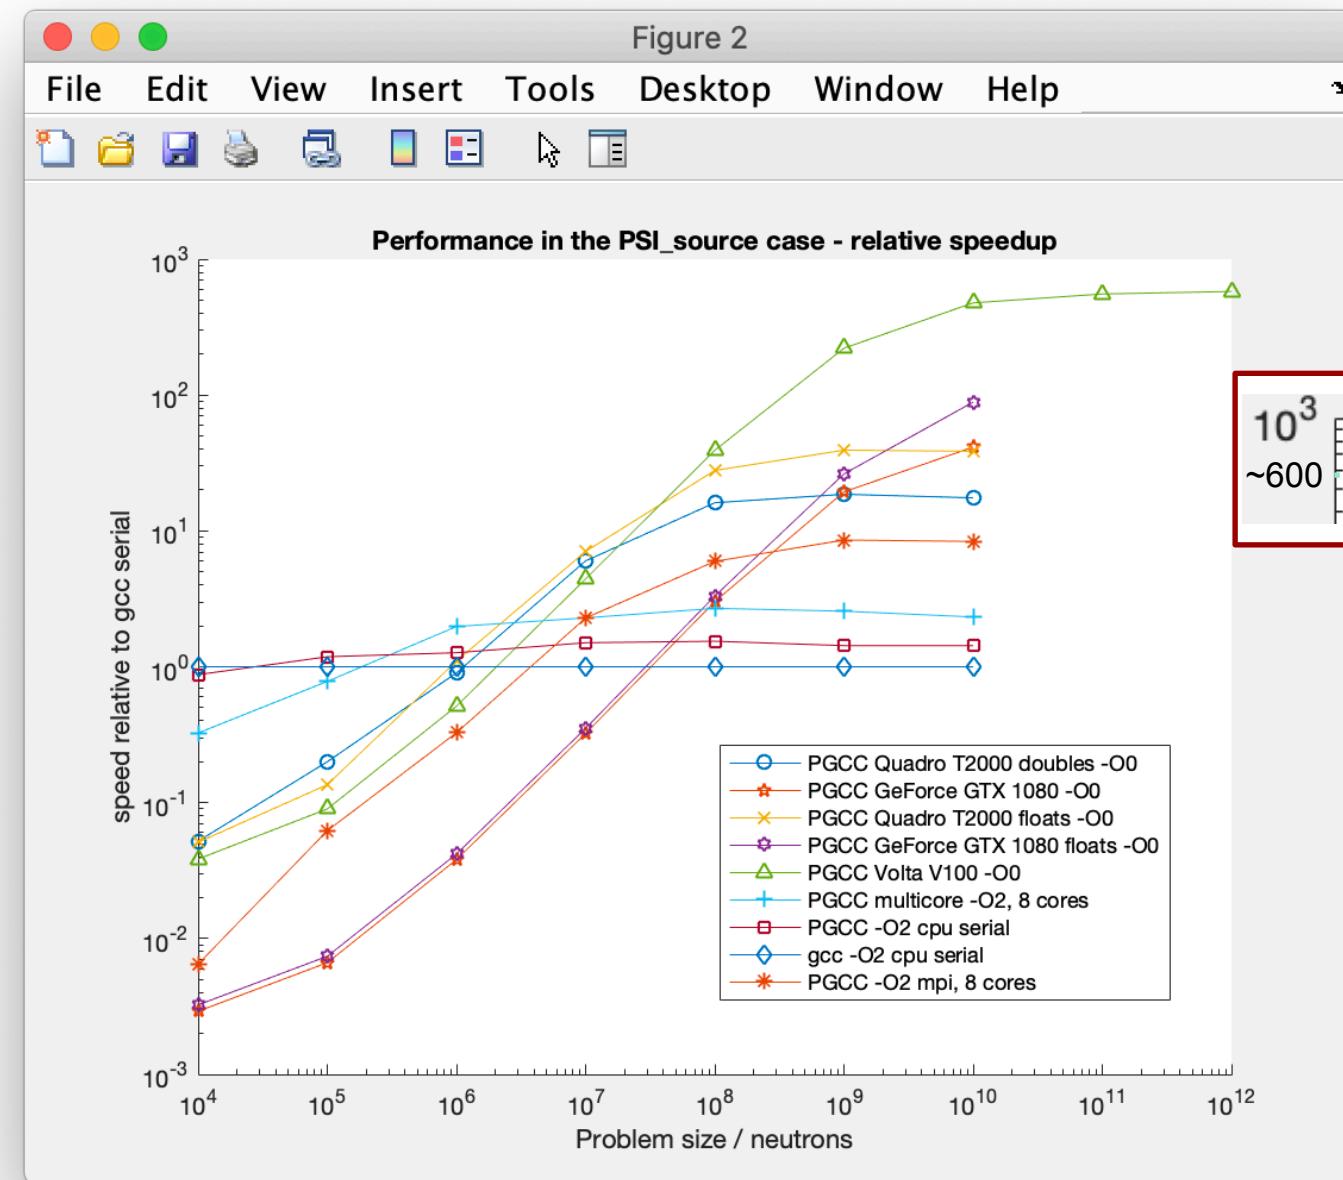
2020 1st Corona lockdown  
P. Willendrup & E. Knudsen  
continue work on comp and cogen

February 2020:  
**First release**  
**McStas 3.0beta**  
with GPU  
support was  
**released**  
to the public

# McStas heading for the GPU... November 2019 - first good look at performance.

**Idealised instrument**  
with source and monitor  
only - i.e. without any  
use of the ABSORB  
macro.

(Likely a good indication  
of maximal speedup  
achievable.)



## Speedup

Looks like a factor of ~600



V100 execution speedups  
renormalised to wall-  
clock of single-core  
gcc standard simulation,

**V100 run is  
600 times faster  
than a single-  
core CPU run**

# McStas heading for the GPU... first benchmarking numbers from November 2019

9 instruments fully ported, also realistic ones like PSI\_DMC

(Aug 2020: 99 instrs)

10-core MPI run,  
 $1\text{e}9$  in 200 secs

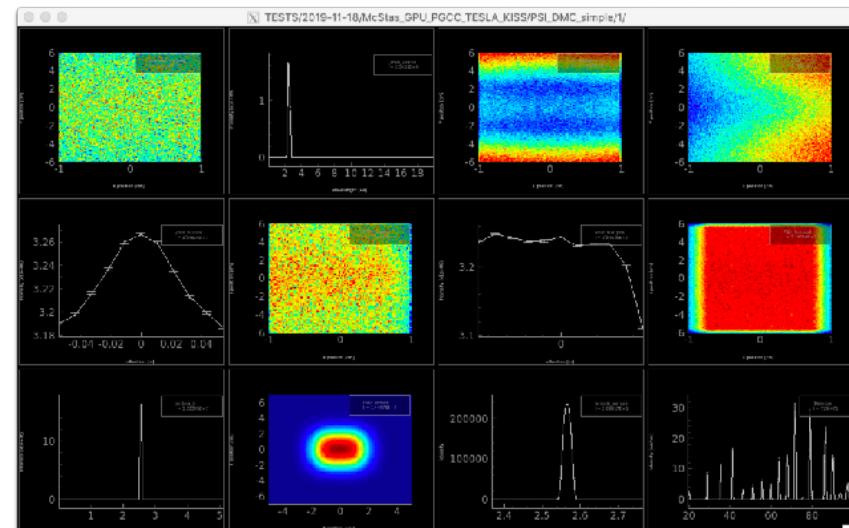
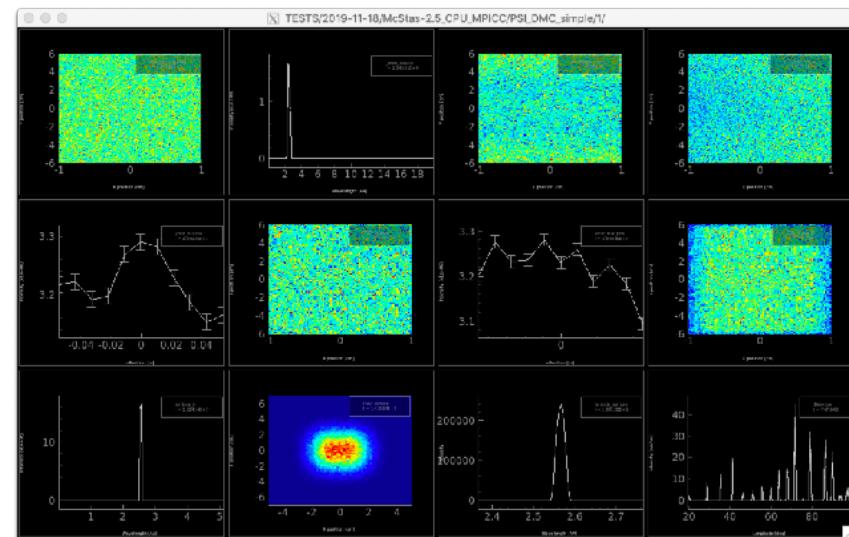


(1-core run,  
 $1\text{e}9$  would be  
2000 secs)

VS.

~ i.e. 2 orders of magnitude wrt. a single, modern CPU core

- If problem has the right size / complexity, GPU via OpenACC is great!



# McStas 2.x -> McStas 3.0 main differences

- **Rewritten / streamlined simplified code-generator** with
  - Much **less generated code**
  - **improved compile time and compiler optimizations**, esp. for large instrs
  - Much **less invasive use of #define**
  - **Component sections -> functions** rather than #define / #undef
  - Much **less global variables**, instrument, component and particle reworked to be **structures**

Advantage of 3.0 also on CPU



- Use of **#pragma acc ...** in lots of places (**put in place by cogen** where possible)

**OpenACC**

- **New random number generator** implemented
  - We couldn't easily port our legacy Mersenne Twister
  - Experimenting with curand showed huge overhead for our relative small number of random numbers  
(we have hundreds or thousands of random numbers, not billions)
- Complete change to **dynamic** monitor-arrays

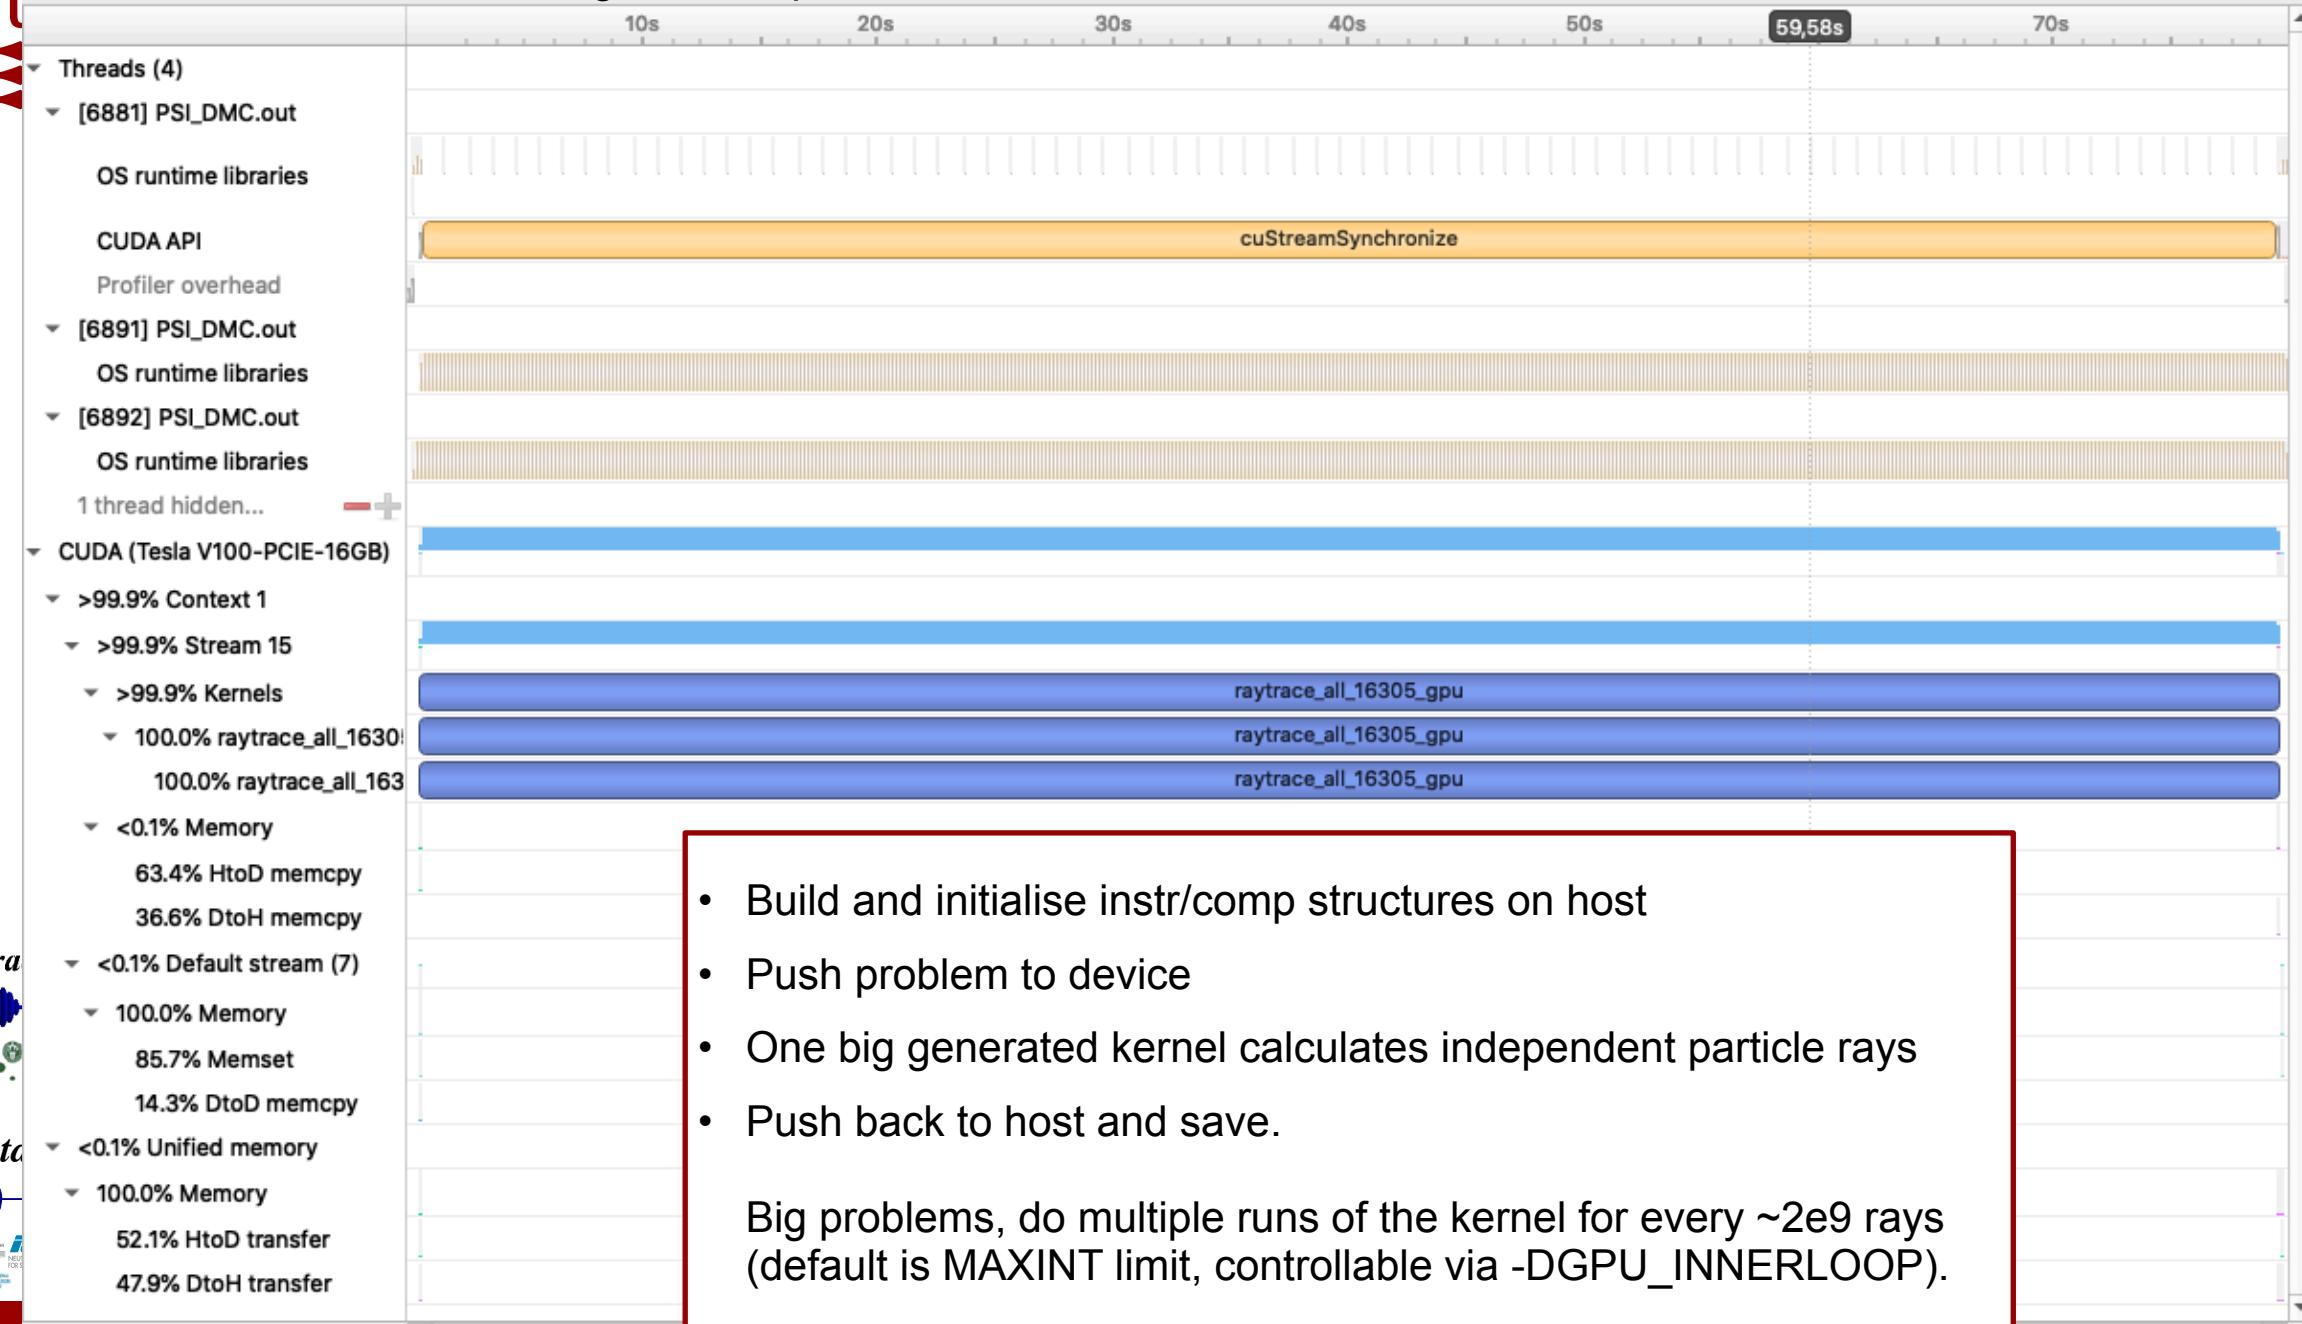
# Anatomy of a McStas GPU run (\*)

- Init, geometry, files etc. read on CPU
  - MPI if needed
- Memory-structures
  - Built on CPU
  - Marked for transfer to GPU (`#pragma acc declare create etc.`)
  - Initialised and synced across
  - Trace-loop is a `#pragma acc parallel loop`
    - Calculation performed entirely on GPU
    - Component structs (incl. e.g. monitor-arrays) synced across
- Finally and Save runs on CPU
  - MPI merge if needed



No printf etc. available  
on GPU, automatically  
suppressed by `#defines`

(\* Alternative layout via FUNNEL mode,  
see next 2 slides)



- Build and initialise instr/comp structures on host
- Push problem to device
- One big generated kernel calculates independent particle rays
- Push back to host and save.

Big problems, do multiple runs of the kernel for every ~2e9 rays  
(default is MAXINT limit, controllable via -DGPU\_INNERLOOP).

# FUNNEL mode

- Activated **explicitly** using -DFUNNEL or **implicitly** using CPUCOMPONENT in instrument or NOACC in comp header

- Has N kernels / calculation zones instead of one

1. Separation at SPLIT

2. Separation if CPUCOMPONENT in instrument file  
( CPUCOMPONENT A=Comp(vars=pars...) )

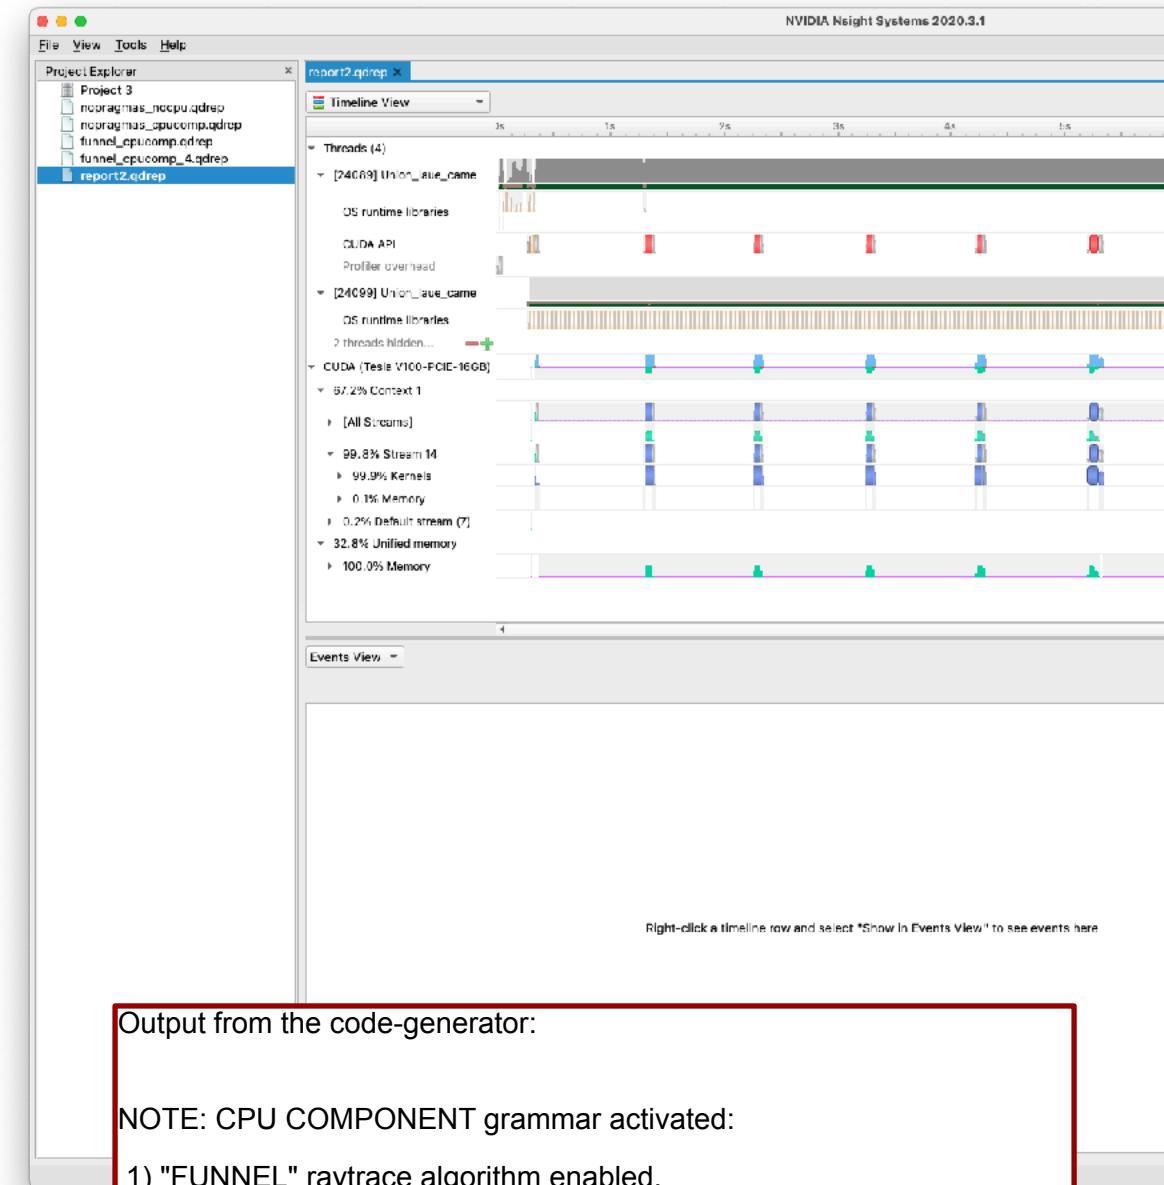
3. Separation if a component has NOACC in the header  
(See e.g. Multilayer\_sample, Union\_master )

- Each of these “calculation zones” is finalised before the next one initiated.

- Example:

Union: Instrument up to Union\_master can be GPU, then CPU, then GPU again

- Can be as slow as single cpu...
- Copying back and forth to/from GPU is costly...



Output from the code-generator:

NOTE: CPU COMPONENT grammar activated:

- 1) "FUNNEL" raytrace algorithm enabled.
- 2) Any SPLIT's are dynamically allocated based on available buffer size.

- Illustration, simple instr with
- Instr vars and “flag”
- Arm
- Source
- Slit
- PSD

```

/* example_v25.instr */
/*
 * %Example: example.instr dummy=0 Detector: detector_I=345.995
 */
DEFINE INSTRUMENT Minimal(dummy=0)

DECLARE ^{
    double constant=2;
    double two_x_dummy;
    double flag;
}

INITIALIZE ^{
    two_x_dummy=2*dummy;
}

TRACE

COMPONENT arm = Arm()
AT (0, 0, 0) ABSOLUTE
EXTEND ^{
    flag=0;
}

COMPONENT source = Source_simple(
    radius = 0.02,
    dist = 3,
    focus_xw = 0.01,
    focus_yh = 0.01,
    lambda0 = 6.0,
    dlambda = 0.05,
    flux = 1e8)
AT (0, 0, 0) RELATIVE arm

COMPONENT coll2 = Slit(
    radius = 0.01)
AT (0, 0, 6) RELATIVE arm
EXTEND ^{
    flag=SCATTERED;
}

COMPONENT detector = PSD_monitor(
    nx = 128,
    ny = 128,
    filename = "PSD.dat",
    xmin = -0.1,
    xmax = 0.1,
    ymin = -0.1,
    ymax = 0.1)
AT (0, 0, 9.01) RELATIVE arm

END

```

```

/* example_v30.instr */
/*
 * %Example: example.instr dummy=0 Detector: detector_I=345.995
 */
DEFINE INSTRUMENT Minimal(dummy=0)

DECLARE ^{
    double constant;
    double two_x_dummy;
}

USERVARS ^{
    double flag;
}

INITIALIZE ^{
    constant=2;
    two_x_dummy=2*dummy;
}

TRACE

COMPONENT arm = Arm()
AT (0, 0, 0) ABSOLUTE
EXTEND ^{
    flag=0;
}

COMPONENT source = Source_simple(
    radius = 0.02,
    dist = 3,
    focus_xw = 0.01,
    focus_yh = 0.01,
    lambda0 = 6.0,
    dlambda = 0.05,
    flux = 1e8)
AT (0, 0, 0) RELATIVE arm

COMPONENT coll2 = Slit(
    radius = 0.01)
AT (0, 0, 6) RELATIVE arm
EXTEND ^{
    flag=SCATTERED;
}

COMPONENT detector = PSD_monitor(
    nx = 128,
    ny = 128,
    filename = "PSD.dat",
    xmin = -0.1,
    xmax = 0.1,
    ymin = -0.1,
    ymax = 0.1)
AT (0, 0, 9.01) RELATIVE arm

END

```

# The particle and USERVARS in the instrument (here shown for McStas)

v2.5: Global variables

```
double x, y, z, vx, vy, vz, t, sx, sy, sz, p;      double flag;
```

v3.0: particle struct, including any USERVARS like flag.

```
struct _struct_particle {
    double x,y,z; /* position [m] */
    double vx,vy,vz; /* velocity [m/s] */
    double sx,sy,sz; /* spin [0-1] */
    unsigned long randstate[7];
    double t, p; /* time, event weight */
    long long _uid; /* event ID */
    long _index; /* component index where to send this event */
    long _absorbed; /* flag set to TRUE when this event is to be removed/ignored */
    long _scattered; /* flag set to TRUE when this event has interacted with the last component instance */
    long _restore; /* set to true if neutron event must be restored */
    // user variables:
    double flag;
};

typedef struct _struct_particle _class_particle;
```

Can be probed using e.g. Monitor\_nD with user1="flag" which uses the function

```
double particle_getvar(_class_particle *p, char *name, int *suc)
also works with e.g. "x"
```

# The particle and USERVARS in the instrument

v2.5: Global variables

```
double x, y, z, vx, vy, vz, t, sx, sy, sz, p;      double flag;
```

RNG state is a thread-variable contained on the \_particle struct. Was earlier a global state in CPU settings

v3.0: particle struct, including any USERVARS like flag.

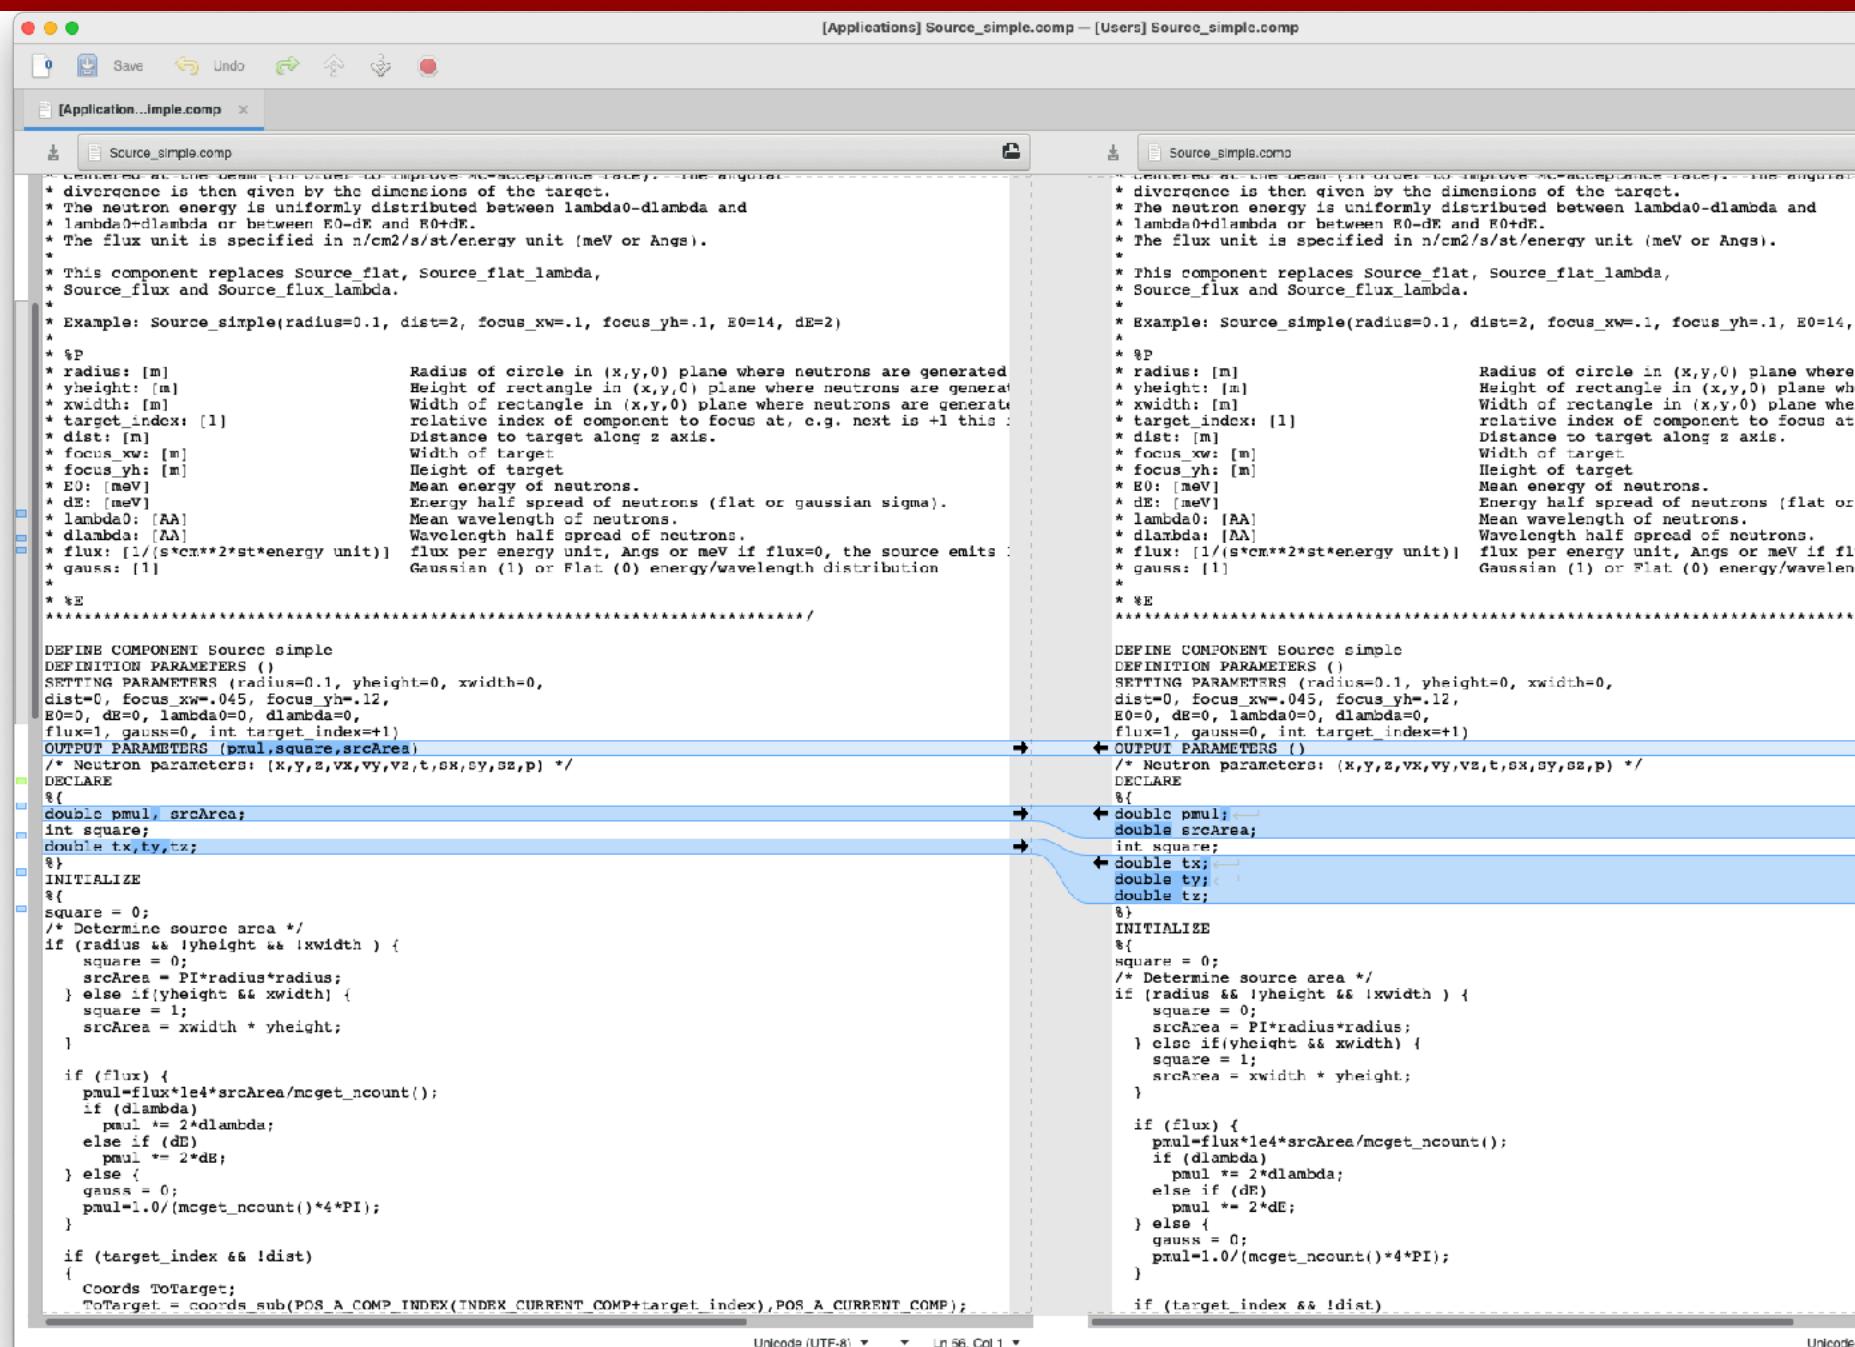
```
struct _struct_particle {
    double x,y,z; /* position [m] */
    double vx,vy,vz; /* velocity [m/s] */
    double sx,sy,sz; /* spin [0-1] */
    unsigned long randstate[7];
    double t, p; /* time, event weight */
    long long _uid; /* event ID */
    long _index; /* component index where to send this event */
    long _absorbed; /* flag set to TRUE when this event is to be removed */
    long _scattered; /* flag set to TRUE when this event has interacted with something */
    long _restore; /* set to true if neutron event must be restored */
    // user variables:
    double flag;
};
typedef struct _struct_particle _class_particle;
```

Side-effect:  
Every function in TRACE that uses random numbers must have \_particle in the footprint

Also:

Particle state data are not global:  
**Don't** use RESTORE\_NEUTRON or RESTORE\_XRAY in TRACE to do a **local** restore, the macro only raises a flag

# Source\_simple minor changes



```

[Applications] Source_simple.comp — [Users] Source_simple.comp

[Application..simple.comp]
Source_simple.comp

Source_simple.comp

* divergence is then given by the dimensions of the target.
* The neutron energy is uniformly distributed between lambda0-dlambda and
* lambda0+dlambda or between E0-dE and E0+dE.
* The flux unit is specified in n/cm2/s/st/energy unit (meV or Angs).
*
* This component replaces Source_flat, Source_flat_lambda,
* Source_flux and Source_flux_lambda.
*
* Example: Source_simple(radius=0.1, dist=2, focus_xw=.1, focus_yh=.1, E0=14, dE=2)
*
* %P
* radius: [m]
* yheight: [m]
* xwidth: [m]
* target_index: [1]
* dist: [m]
* focus_xw: [m]
* focus_yh: [m]
* E0: [meV]
* dE: [meV]
* lambda0: [AA]
* dlambda0: [AA]
* flux: [1/(s*cm**2*st*energy unit)]
* gauss: [1]
Radius of circle in (x,y,0) plane where neutrons are generated
Height of rectangle in (x,y,0) plane where neutrons are generated
Width of rectangle in (x,y,0) plane where neutrons are generated
relative index of component to focus at, e.g. next is +1 this :
Distance to target along z axis.
Width of target
Height of target
Mean energy of neutrons.
Energy half spread of neutrons (flat or gaussian sigma).
Mean wavelength of neutrons.
Wavelength half spread of neutrons.
flux per energy unit, Angs or mev if flux=0, the source emits :
Gaussian (1) or Flat (0) energy/wavelength distribution
*
* #E
*****/



DEFINE COMPONENT Source_simple
DEFINITION PARAMETERS ()
SETTING PARAMETERS (radius=0.1, yheight=0, xwidth=0,
dist=0, focus_xw=.045, focus_yh=.12,
E0=0, dE=0, lambda0=0, dlambda0=0,
flux=1, gauss=0, int target_index+1)
OUTPUT PARAMETERS (paul,square,srcArea)
/* Neutron parameters: (x,y,z,vx,vy,vz,t,sx,sy,sz,p) */
DECLARE
{
    double pmul, srcArea;
    int square;
    double tx,ty,txz;
}
INITIALIZE
{
    square = 0;
    /* Determine source area */
    if (radius && lyheight && lxwidth ) {
        square = 0;
        srcArea = PI*radius*radius;
    } else if(yheight && xwidth) {
        square = 1;
        srcArea = xwidth * yheight;
    }

    if (flux) {
        paul=flux*le4*srcArea/mcget_ncount();
        if (dlambda)
            pmul *= 2*dlambda;
        else if (dE)
            pmul *= 2*dE;
        else {
            gauss = 0;
            paul=1.0/(mcget_ncount()*4*PI);
        }
    }

    if (target_index && !dist)
    {
        Coords ToTarget;
        ToTarget = coords_sub(POS_A_COMP_INDEX(TINDEX_CURRENT_COMP+target_index),POS_A_CURRENT_COMP);
    }
}

* divergence is then given by the dimensions of the target.
* The neutron energy is uniformly distributed between lambda0-dlambda and
* lambda0+dlambda or between E0-dE and E0+dE.
* The flux unit is specified in n/cm2/s/st/energy unit (meV or Angs).
*
* This component replaces Source_flat, Source_flat_lambda,
* Source_flux and Source_flux_lambda.
*
* Example: Source_simple(radius=0.1, dist=2, focus_xw=.1, focus_yh=.1, E0=14, dE=2)
*
* %P
* radius: [m]
* yheight: [m]
* xwidth: [m]
* target_index: [1]
* dist: [m]
* focus_xw: [m]
* focus_yh: [m]
* E0: [meV]
* dE: [meV]
* lambda0: [AA]
* dlambda0: [AA]
* flux: [1/(s*cm**2*st*energy unit)]
* gauss: [1]
Radius of circle in (x,y,0) plane where
Height of rectangle in (x,y,0) plane where
Width of rectangle in (x,y,0) plane where
relative index of component to focus at
Distance to target along z axis.
Width of target
Height of target
Mean energy of neutrons.
Energy half spread of neutrons (flat or
Mean wavelength of neutrons.
Wavelength half spread of neutrons.
flux per energy unit, Angs or mev if flux=0, the source emits :
Gaussian (1) or Flat (0) energy/wavelength distribution
*
* #E
*****/



DEFINE COMPONENT Source_simple
DEFINITION PARAMETERS ()
SETTING PARAMETERS (radius=0.1, yheight=0, xwidth=0,
dist=0, focus_xw=.045, focus_yh=.12,
E0=0, dE=0, lambda0=0, dlambda0=0,
flux=1, gauss=0, int target_index+1)
OUTPUT PARAMETERS ()
/* Neutron parameters: (x,y,z,vx,vy,vz,t,sx,sy,sz,p) */
DECLARE
{
    double pmul;
    double srcArea;
    int square;
    double tx;
    double ty;
    double tz;
}
INITIALIZE
{
    square = 0;
    /* Determine source area */
    if (radius && lyheight && lxwidth ) {
        square = 0;
        srcArea = PI*radius*radius;
    } else if(yheight && xwidth) {
        square = 1;
        srcArea = xwidth * yheight;
    }

    if (flux) {
        pmul=flux*le4*srcArea/mcget_ncount();
        if (dlambda)
            pmul *= 2*dlambda;
        else if (dE)
            pmul *= 2*dE;
        else {
            gauss = 0;
            pmul=1.0/(mcget_ncount()*4*PI);
        }
    }

    if (target_index && !dist)
    {
        Coords ToTarget;
        ToTarget = coords_sub(POS_A_COMP_INDEX(TINDEX_CURRENT_COMP+target_index),POS_A_CURRENT_COMP);
    }
}

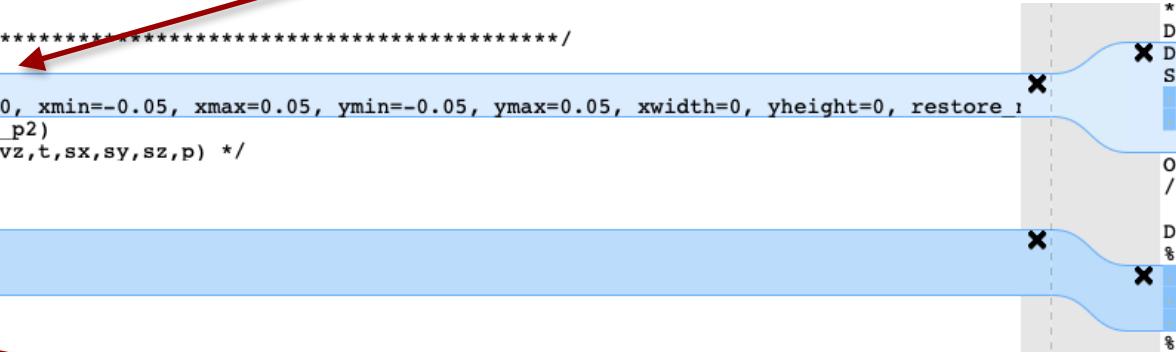
```

# PSD has several changes

## No more DEFINITION PARAMETERS

```
* %E
*****
***** DEFINE COMPONENT PSD_monitor *****
***** DEFINITION PARAMETERS (nx=90, ny=90) *****
***** SETTING PARAMETERS (string filename=0, xmin=-0.05, xmax=0.05, ymin=-0.05, ymax=0.05, xwidth=0, yheight=0, restore_neutron=0, int nowritefile=0) *****
***** OUTPUT PARAMETERS (PSD_N, PSD_p, PSD_p2)
/* Neutron parameters: (x,y,z,vx,vy,vz,t,sx,sy,sz,p) */

DECLARE
%{
double PSD_N[nx][ny];
double PSD_p[nx][ny];
double PSD_p2[nx][ny];
%}
INITIALIZE
%{
```



```
***** DEFINE COMPONENT PSD_monitor *****
***** X DEFINITION PARAMETERS () *****
***** X SETTING PARAMETERS (nx=90, ny=90, string filename=0,
***** X xmin=-0.05, xmax=0.05, ymin=-0.05, ymax=0.05, xwidth=0, yheight=0,
***** X restore_neutron=0, int nowritefile=0) *****
***** X OUTPUT PARAMETERS (PSD_N, PSD_p, PSD_p2)
/* Neutron parameters: (x,y,z,vx,vy,vz,t,sx,sy,sz,p) */

DECLARE
%{
X DArray2d PSD_N;
X DArray2d PSD_p;
X DArray2d PSD_p2;
%}

INITIALIZE
%{
if (xwidth > 0) { xmax = xwidth/2; xmin = -xmax; }
if (yheight > 0) { ymax = yheight/2; ymin = -ymax; }

if ((xmin >= xmax) || (ymin >= ymax)){
    printf("PSD_monitor: %s: Null detection area !\n"
        "ERROR          (xwidth,yheight,xmin,xmax,ymin,ymax). Exiting",
        NAME_CURRENT_COMP);
    exit(0);
}

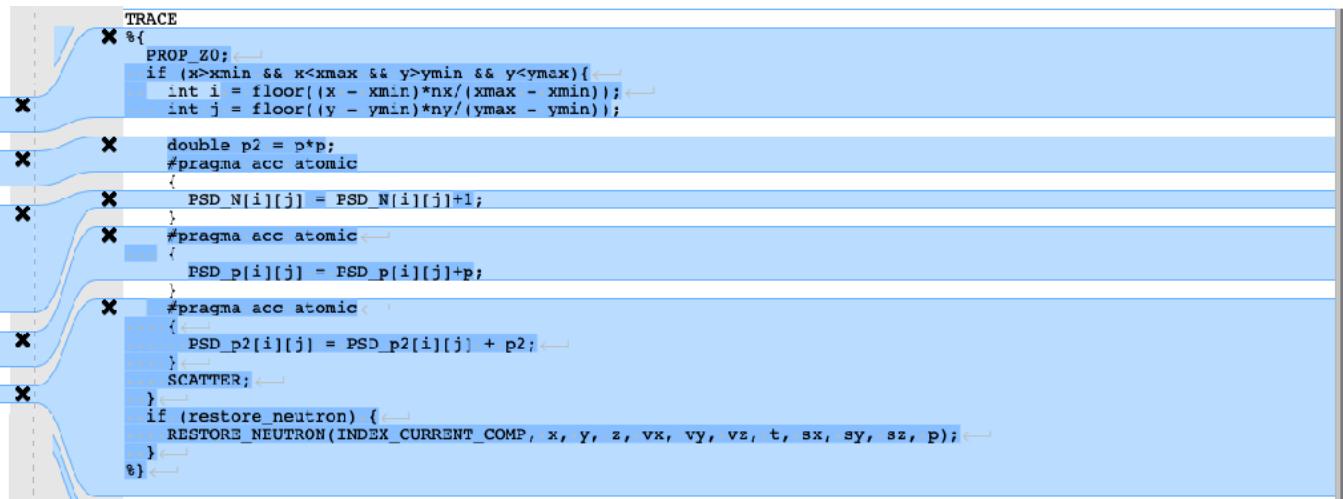
PSD_N = create_darr2d(nx, ny);
PSD_p = create_darr2d(nx, ny);
PSD_p2 = create_darr2d(nx, ny);

%}
```

Use of new DArray2d for dynamic allocation

# PSD lots of changes

```
TRACE
  {
    int i,j;
    PROP_Z0;
    if (x>xmin && x<xmax && y>ymin && y<ymax)
    {
      i = floor((x - xmin)*nx/(xmax - xmin));
      j = floor((y - ymin)*ny/(ymax - ymin));
      PSD_N[i][j]++;
      PSD_p[i][j] += p;
      PSD_p2[i][j] += p*p;
      SCATTER;
    }
    if (restore_neutron) {
      RESTORE_NEUTRON(INDEX_CURRENT_COMP, x, y, z, vx, vy, vz, t, sx, sy, sz, p);
    }
  }
```



```
PROP_Z0;
if (x>xmin && x<xmax && y>ymin && y<ymax){
  int i = floor((x - xmin)*nx/(xmax - xmin));
  int j = floor((y - ymin)*ny/(ymax - ymin));
```

```
double p2 = p*p;
#pragma acc atomic
{
  PSD_N[i][j] = PSD_N[i][j]+1;
```

```
}
```

```
#pragma acc atomic
```

```
{
  PSD_p[i][j] = PSD_p[i][j]+p;
```

```
}
```

```
#pragma acc atomic
```

```
{
  PSD_p2[i][j] = PSD_p2[i][j] + p2;
```

```
}
```

```
SCATTER;
```

Enabling atomic writes on the detector arrays



Peter Willendrup, DTU Physics and ESS DMSC

16

# Samples required a good deal of (detective) work...

Key issue: (mis-) Use of / component DECLARE variables in TRACE for storing particle-dependent information, e.g. reflection list in PowderN etc. must be avoided.

Solutions:

1) Make local thread-variables, e.g. →

2) Where meaningful, one could make atomic sections ala the monitors

Side-effect of thread-local vars:

Next neutron(s) in a SPLIT are no longer aware of e.g. available powder lines.

Potential, future solution: Mechanism to inject comp-specific code in the \_particle

```
// Variables calculated within thread for thread purpose only
char type = '\0';
int itype = 0;
double d_phi_thread = d_phi;
// These ones are injected back to struct at the end of TRACE in non-OpenACC case
int nb_reuses = line_info.nb_reuses;
int nb_refl = line_info.nb_refl;
int nb_refl_count = line_info.nb_refl_count;
double vcache = line_info.v;
double Nq = line_info.Nq;
double v_min = line_info.v_min;
double v_max = line_info.v_max;
double lfree = line_info.lfree;
double neutron_passed = line_info.neutron_passed;
long xs_compute = line_info.xs_compute;
long xs_reuse = line_info.xs_reuse;
long xs_calls = line_info.xs_calls;
int flag_warning = line_info.flag_warning;
double dq = line_info.dq;

#ifndef OPENACC
#ifndef USE_OFF
off_struct thread_offdata = offdata;
#endif
#else
#define thread_offdata offdata
#endif
```

# New monitor-tools for debugging use

- **Event\_monitor\_simple(nevents=1e6)**
  - basic non-Monitor\_nD event monitor. Writes a “log” file, independent from detector\_out macros
- **Flex\_monitor\_1D , Flex\_monitor\_2D, Flex\_monitor\_3D,**
  - simple 1/2/3D “uservar” monitors tapping into the instrument USERVARS ala Monitor\_nD
- Useful for **debugging even component** internals:  
On mccode-3, if same ncount, same seed, same level of MPI parallelisation, the output should be **identical** on CPU and GPU

**Each component will correspond to a set of functions. Trace is a GPU'ified function...**

**+ particle-loop and logic around, also running on GPU.**

**Init and finalisation codes run purely CPU.**

```
#pragma acc routine seq
_class_Source_simple *_class_Source_simple_trace(_class_Source_simple *_comp
, _class_particle *_particle) {
ABSORBED=SCATTERED=RESTORE=0;

#define radius (_comp->_parameters.radius)
#define yheight (_comp->_parameters.yheight)
#define xwidth (_comp->_parameters.xwidth)
#define dist (_comp->_parameters.dist)
#define focus_xw (_comp->_parameters.focus_xw)
#define focus_yh (_comp->_parameters.focus_yh)
#define E0 (_comp->_parameters.E0)
#define dE (_comp->_parameters.dE)
#define lambda0 (_comp->_parameters.lambda0)
#define dlambd a (_comp->_parameters.dlambd a)
#define flux (_comp->_parameters.flux)
#define gauss (_comp->_parameters.gauss)
#define target_index (_comp->_parameters.target_index)
#define pmul (_comp->_parameters.pmul)
#define square (_comp->_parameters.square)
#define srcArea (_comp->_parameters.srcArea)
#define tx (_comp->_parameters.tx)
#define ty (_comp->_parameters.ty)
#define tz (_comp->_parameters.tz)
SIG_MESSAGE("[_source_trace] component source=Source_simple() TRACE [Source_simple.comp:127]");
double chi,E,lambda,v,r, xf, yf, rf, dx, dy, pdir;
t=0;
z=0;

if (square == 1) {
    x = xwidth * (rand01() - 0.5);
    y = yheight * (rand01() - 0.5);
} else {
    chi=2*PI*rand01();
    r=sqrt(rand01())*radius;
    x=r*cos(chi);
    y=r*sin(chi);
}
randvec_target_rect_real(&xf, &yf, &rf, &pdir,
                         tx, ty, tz, focus_xw, focus_yh, ROT_A_CURRENT_COMP, x, y, z, 2);
.... etc
```

“Scatter-gather” approach not far from what we do in MPI settings, i.e. :

GPU case:

N particles are calculated in parallel in N GPU threads. (Leave to OpenACC/device how many actually are running at one time)

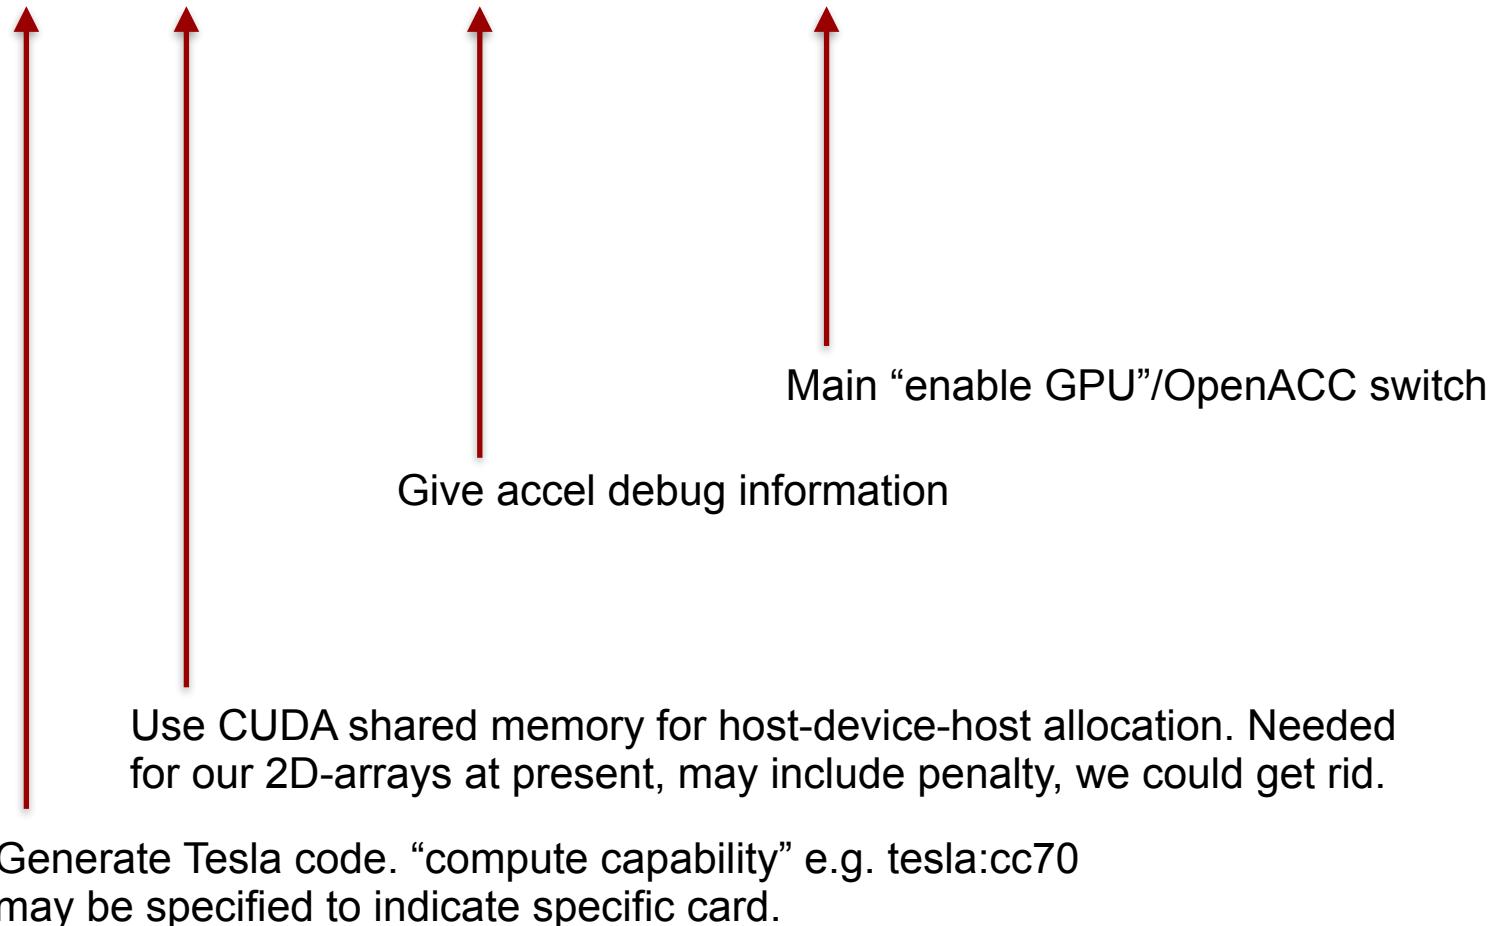
CPU case:

N particles are calculated in M serial chunks over M processors.

Contains component trace section

# Compiler settings used for GPU:

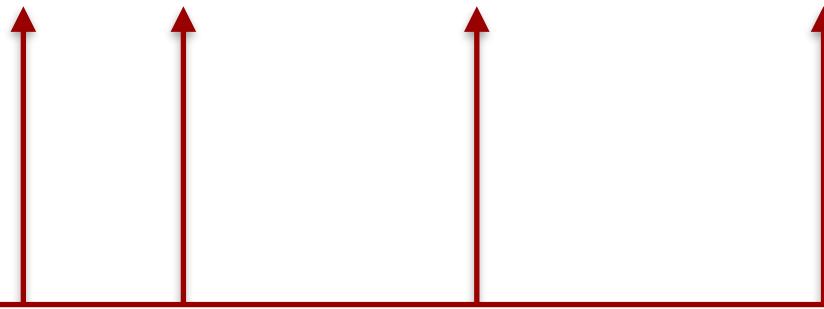
```
nvc -ta=tesla,managed -Minfo=accel -DOPENACC
```



( McStas 3.0 mcrun is preconfigured on Linux - excluding -Minfo=accel, simply use mcrun --openacc when compiling, can also combine with e.g. --mpi=N )

# Compiler settings used for GPU:

```
nvc -ta=tesla,managed -Minfo=accel -DOPENACC
```



Defaults for “GPU neutron loops”:

1) non-funnelled

**GPU\_INNERLOOP=2147483647**

2) funnelled

**GPU\_FUNNEL\_INNERLOOP=1024\*1024**

For CPU/threading, use below settings, with -DMULTICORE e.g. printf's are not nullified in TRACE

```
nvc -ta=multicore -Minfo=accel -DOPENACC ( -DMULTICORE )
```

# What doesn't work

- Function pointers are not available on GPU
  - Solutions:
    - Code around if possible (integration routine pr. specific function to be integrated...)
    - Mark the component NOACC
  - Variadic functions are not available on GPU
  - Anonymous structs as comp pars are not available on GPU
    - Unfold into comp struct
  - User-defined fieldfunctions for polarisation had to be abandoned
    - No solution yet, may become handled via grammar
  - External libs generally can not be used in TRACE ("#pragma...." hard to add on 3rd party codes)
    - Handle in INIT / FINALLY (MCPL)
    - NOACC (GSL etc.)
  - Union master is for now NOACC, will eventually become supported on GPU



# Porting an instrument to 3.0 and GPU

- **Instrument-level variables** that are to become particle-dependent “**flags**”, e.g. for use in EXTEND WHEN must be put in the new section **USERVARS %{ double flag; %}**
- Use of **instrument input pars** in extend and WHEN must use **INSTRUMENT\_GETPAR(varname)**
- **Non-flag instrument vars** to be used during TRACE / EXTEND / WHEN must be injected via **#pragma acc declare create(var)** and **#pragma acc update device(var)**
- **Declare-functions** to be used in **trace** (e.g. in an EXTEND) must have **#pragma acc routine**



# Common error-messages:

- At compile-time, nvc is quite informative if using -Minfo:accel

PGC-S-0155-Invalid atomic expression

PGC-S-0000-Internal compiler error. Error: Detected unexpected atomic store opcode.

PGC-S-0155-External variables used in acc routine need to be in #pragma acc create() - flag

...

- At runtime, this indicates a GPU segfault

Failing in Thread:1

call to cuMemcpyDtoHAsync returned error 700: Illegal address during kernel execution

- Often a symptom of “illegal access”, colliding memory, something isn’t thread-safe...

# Porting a comp to GPU

- All pars must be setting parameters.

New types: string a="none" and vector b ( either ={1,2,3,4} static init or via pointer.)

- All function-declarations must be moved to SHARE

- DECLARE must **only** contain variable declarations. All initialization resides in INITIALIZE

- If the comp uses external libs either

- Avoid use in component TRACE (e.g. MCPL\_input and output, handled in INIT/FINALLY)
  - Use NOACC keyword (e.g. Multilayer\_sample use of GSL) - implies FUNNEL mode



- Add #pragma acc routine to functions to execute in TRACE

- Functions that call rand01() and friends must include \_class\_particle \*\_particle in footprint.  
(rand01() etc. are macros that carry thread-seed on \_particle)s

- Generally, don't store ANY particle-derived vars on comp struct, make local TRACE vars instead.
  - Exception: Monitors, handle arrays in #pragma acc atomic clauses

- Don't use RESTORE\_XRAY/NEUTRON in TRACE to do a local restore, the macro only raises a flag

# Highlights of comps that work differently

- Monitor\_nD
  - uservars are strings user1="flag", they use \_particle\_getvar to access instrument USERVARS
- MCPL\_input and MCPL\_output
  - do most of their work in INIT/FINALLY - buffers transferred for TRACE use
- PowderN + Single\_crystal + Isotropic\_sqw
  - don't check for "same particle as before"
  - in SPLIT cases, no particle state info is kept
  - (we could potentially use \_particle and "USERVARS" injected from the comps...)

# Conclusions

- **It really does work nicely!**
- **Code changes** much **less invasive** than envisioned!
- It often gives a speedup of **1-2 orders** of magnitude over 1 cpu
- **Most things work**  
(we have workarounds or solutions in the pipe for the rest)
- **Documentation** comes in the form of the released code + this set of slides...
- McXtrace 3.0 and McStas 3.1 fully support GPU but **not fully “optimised” performance-wise**, we will try to go to another Hackathon
- **Union** needs a dedicated **Hackathon**
- **Compilation** with GCC 10 offloading support achieved May 2021 - but produces 0 on detectors...  
- hope for better GCC support and non-NVIDIA cards in 1-2 years



# The team, Nvidia mentors and Hackathon hosts :-)



Guido Juckeland



Sebastian von Alfthan



Vishal Metha



Christian Hundt



Alexey Romanenko

